ABA Panel — Working Notes

Started: February 14, 2026 **Process:** Review synthesis topics and source docs; capture reactions, cuts, slide ideas, and phrasing.

Block 0: How LLMs Work (Law Clerk Analogy)

Source to review: `Legal_AI_Vision_Claude.docx`, Block 0 script in `02-panel-prep.md`

Reactions / edits to the script:

Keep / cut / add:

Slide ideas:

I. The Regulatory Landscape (427 Orders)

Sources to review: `ropes_gray_tracker_data.md`, `local_rules/` directory, `AI Annual Report 2025.pdf` (Part 161 proposal), `Illinois.pdf`

Reactions:

Key points I want to make:

Points I'm cutting or leaving to other panelists:

Slide ideas:

- **"The Spectrum" table** — one slide showing the full range of judicial approaches:

Approach	Examples	Trend
Outright ban	Judge Boyko, N.D. Ohio (carve-out for Westlaw/Lexis)	Rare and shrinking group. Bans generally adopted in 2023/2024 timeframe.
Permissive, no disclosure	Illinois Supreme Court statewide policy (eff. Jan 1, 2025): AI use "should not be discouraged, and is authorized," so long as it complied with legal and ethical standards	Minority position but influential
Brief-only disclosure	N.D. Tex. LR 7.2(f) — first-page disclosure; omission = negative certification	Early model (Sept 2024)
All-court filings disclosure	Judge Padin, D.N.J.; W.D. Okla. Bankruptcy (one of earliest, July 2023)	Most common approach

Approach	Examples	Trend
Filings + correspondence	Judge Pratter, E.D. Pa. — specifically extends to correspondence with court	Broadest disclosure scope
Certification + record-keeping	Judge Martinez-Olguin, N.D. Cal. — must retain all prompts/queries "in the event those records become relevant"	Most demanding compliance burden
Preparation vs. drafting distinction	Judge Broderick, S.D.N.Y. — research use needs no disclosure; drafting requires certification; M.J. Kang, N.D. Cal. — exempts "traditional legal research, word processing, spellchecking, grammar checking, or formatting software tools."	Most nuanced; gaining influence

- **"Trends" slide** — companion to the Spectrum table:
  - Over 200 orders added in H2 2025 alone; pace intensifying into Jan-Feb 2026. All 50 states now represented.
  - Bans are fading — Ropes Gray map shows essentially no states in "prohibits" (red); even Judge Boyko's ban carves out traditional research platforms
  - Disclosure is the floor — virtually every jurisdiction expects some form of transparency about AI use
  - Certification is emerging as the next layer — signed attestation of human verification, especially in federal courts
  - More focus on professional responsibility rules — courts increasingly invoke Rules 1.1, 3.3, and 11 rather than creating novel AI-specific obligations
- **"Contrasting State Approaches" slide** — three poles:
  - **Illinois** Supreme Court (eff. Jan 2025): AI use is "authorized," disclosure "should not be required in a pleading," and accountability flows through existing professional conduct rules.
  - **New York** Advisory Committee (Dec 2025): proposed Part 161 model rule opposes blanket bans but requires meaningful disclosure. Notably recommends against certification requirements, favoring existing 22 NYCRR 130-1.1 framework (papers contain no false or fictitious material).
  - **California** SB 574 (passed Senate Jan 19, 2026): requires attorneys to take "reasonable steps" to verify all AI-generated materials, correct hallucinations, and remove biased content. Prohibits inputting confidential client information into public AI tools. Bars arbitrators from delegating decisions to AI.

II. The Hallucination Crisis

Sources to review: `Charlotin-hallucination_cases.csv`, `Noland`, `Nuvola`, `Flycatcher`, `Cassata`, Jenner & Block/Kjoller

II-A. The Scale (data / statistics)

Reactions:

Which numbers to use (less is more):

Slide ideas:

- **"The Scale" slide** — the epidemic in six lines:
  - **851 documented cases** as of late January 2026, spanning **30 countries** [Charlotin database]
  - **716 cases in 2025 alone** — a 10x increase over all prior years combined
  - **130 in December 2025 alone** — more than 4 per day
  - **40% involve lawyers** (not just pro se litigants) — 342 cases
  - **9 cases involve judges** who used AI in decisions
  - A 2024 Stanford study found that even premium legal research platforms hallucinate — Westlaw AI at 33%, Lexis AI at 17%. Those rates are likely improving as platforms build in verification tools, but subtler forms of hallucination — misrepresented holdings, jurisdictional blurring — are harder to catch automatically.

II-B. Fabricated Citations / Escalating Consequences

Reactions after reading the cases:

Which cases to tell as stories vs. just cite:

- **"Escalating Consequences" slide** — sanctions table:

Case	Court	Sanction	Significance
Mata v. Avianca (2023)	S.D.N.Y.	$5,000 fine	The original; established the paradigm
Noland v. Land of the Free (Sept 2025)	Cal. App.	$10,000 + State Bar referral	21 of 23 quotations fabricated; attorney used ChatGPT, Claude, Gemini, and Grok sequentially without reading the briefs
Nuvola (Nov 2025)	MN District	$1,000 + OLPR referral + 5 mandatory CLE presentations	Creative sanctions; also admonished opposing counsel for not catching fake cites
Deutsche Bank v. LeTennier (Jan 2026)	NY App. Div. 3d	$5,000 against attorney	First NY appellate sanction for GenAI misuse (at least 23 fabricated cites)
Flycatcher/Feldman (Feb 2026)	S.D.N.Y.	Default judgment	Terminal sanction — most severe to date; attorney submitted fake cites across three separate filings despite explicit court warnings

Case	Court	Sanction	Significance
LaBatt USA v. Friends Bev. (Dec 2025)	NY Comm Div	Striking opposition to SJ; motion treated as unopposed	Court awarded $1.1 million on summary judgment after striking opposition for having twice misused AI in the case
Kjoller v. Superior Court (2026)	Cal. Supreme Court	OSC + referee investigation	DA cited 8 cases — 3 fabricated, 3 misrepresented — then blamed "scrivener's errors"; Supreme Court unanimously directed formal sanctions inquiry via CCP §§ 638-640 referee process

II-B2. The Detection Problem

- **"The Detection Problem" slide:**
  - Courts are developing informal detection methods — stylistic analysis, comparing a lawyer's other filings, circumstantial evidence like a firm's recent AI tool acquisition.
  - But there's no reliable technical solution to detection, and there never will be.
  - **Why detection can't solve this:** AI text is generated by predicting the most likely next word based on patterns learned from human writing. There's no hidden signature — just words that follow the same statistical patterns humans use. And it's an arms race: as detectors learn to identify AI patterns, the models are retrained to avoid them. The generators will always be a step ahead.
  - **So the right question isn't "was AI used?" — it's "is the content accurate?"** Which is why verification matters more than disclosure.

II-C. Beyond Fake Citations: Subtler Forms of Hallucination

- **"Beyond Fake Citations" slide:**
  - **Blurring / Cross-jurisdictional contamination:** When you search Westlaw for New York cases, it applies a filter — the database cannot return a California case. An LLM doesn't work that way. It was trained on the whole library, blended into statistical patterns with no jurisdictional labels. When you say "only consider New York law," you're asking the model to try not to use part of what it learned. The result: answers that may apply the wrong jurisdiction's rule or subtly blend majority and minority approaches.
  - **Disregarding your sources:** You give the LLM a document to analyze, but it has a technical problem reading it — so it guesses what a "typical" document would say and gives you a confident, well-written analysis of something it never read. I've actually seen LLMs do this — you can watch the thinking traces with some models, and I've literally seen the LLM "think": "I cannot access the uploaded file. I'll assume it contains the typical elements..."
  - **Why these are arguably worse than fake citations:** Fake citations are easy to catch — you look them up. Blurring and fabricated analysis look right. The case exists. The rule sounds correct. But it's the wrong jurisdiction's rule or an analysis of a document the AI never actually read.

Still to decide: Do I need the N.Z. v. Fenix opinion to support the blurring point?

II-D. Disregarding Specified Sources

My own experiences to share:

How much personal anecdote is right here:

II-E. The Cover-Up Pattern (Kjoller / Hayes)

Reactions after reading Jenner & Block piece:

How to frame this for a business law audience:

- **"The Cover-Up Is Worse Than the Crime" slide:**
  - **Cassata** (NY Sup. Ct., Jan 2026): Attorney denied using AI. Court found it "more likely than not" based on circumstantial evidence — firm had just acquired AI-enhanced Westlaw, supervising attorney admitted he didn't know "how to use any AI stuff." The brief also plagiarized from a filing in another case that had itself been condemned for AI fabrications.
  - **Kjoller** (Cal. Supreme Court, 2026): After opposing counsel discovered fabricated citations, the DA tried to minimize — first calling to say she'd been rushing, then filing a brief characterizing wholesale fabrication as "scrivener's errors." The Supreme Court unanimously directed a formal sanctions inquiry, transforming a correctable mistake into a potential career-ending investigation.
  - **Hayes** (E.D. Cal.): Attorney blamed "hasty" drafting. Court ordered sanctions notice sent to every state bar where the attorney was licensed and every judge in the district.
  - **The emerging rule:** Acknowledge and correct = manageable consequences. Deflect, deny, or minimize = investigations, public embarrassment, escalating sanctions.
  - **Note on Kjoller:** This was a criminal case — a person's life and liberty were at stake, not just a contractual dispute.

- **Talking point after the cover-up slide (direct address to audience):**

  > "The cover-up cases involve people who compounded the problem with dishonesty — denying AI use, calling fabrications 'scrivener's errors,' lying to courts. Those aren't the people in this room. You're here at ABA, attending CLE, maintaining your technological competence. You're not going to submit fake citations and then cover it up.
  >
  > But here's what I want you to take away: this problem isn't limited to bad actors or lazy lawyers. In my own experience on the bench, I've been genuinely surprised by good lawyers — experienced, ethical practitioners — who didn't fully understand what they were dealing with and got into trouble. It's not always who you think. The technology is new, levels of sophistication vary enormously, and the tools are designed to sound confident whether they're right or wrong. That's why understanding how these systems actually work matters — which is where we started today."

II-F. Premium Platform Hallucination Rates

Reactions:

How to deliver the 17% / 33% / 58-88% numbers:

III. Professional Responsibility

Sources to review: `ABA BLS slides`, `Comfed_AI_Jan2026.pdf` (Opinion 512), `AI Annual Report 2025.pdf`

III-A. Model Rules Quick-Run

Which rules to hit vs. skip:

Reactions:

III-B. Verification Duty (it doesn't matter whether AI was used)

Reactions:

Do I need Short Load Concrete / Shayan v. Shakib? Notes:

III-C. Supervisory Responsibility

Reactions after reading Cassata closely:

The Nuvola opposing-counsel-duty point — use it or not:

III-D. Duty of Technological Competence

Reactions:

How to frame "cuts both ways" for this audience:

III-E. Legislative Turn (SB 574, EO 231, NY bills)

Sources to review: `AI Legislation Tracker _2026_.pdf`, Jenner & Block/Kjoller

Reactions:

What the business law audience needs to hear vs. what's inside baseball:

III-F. Disclosure Landscape

Reactions:

My own view on disclosure vs. certification:

IV. Judicial Use of AI

Sources to review: `Sedona/Dixon`, `CofC handout` (pp. 1-20), `Comfed`, `a.i.-policy.pdf`, `Fordham handout`

IV-A. What Judges Are Doing

Reactions:

What I'm comfortable sharing about my own practice:

IV-B. Ethical Considerations (Rule 2.9(C), bias, confidentiality)

Reactions:

Open questions I want to raise with the audience:

IV-C. Sedona Guidelines

Reactions:

Which use cases to highlight:

IV-D. My Vision: Verifiable Assistance

Source to review: `Legal_AI_Vision_Claude.docx`

What I want to say here — in my own words:

How much to share about the NCE system:

The closing line — draft and refine:

Cross-Cutting Notes

Connections between topics I noticed:

Things I want to say that don't fit neatly into one section:

Points I expect other panelists to cover (so I can defer):

Anticipated audience questions and my answers:

1.
2.
3.
4.
5.

Slide Tracker

As slide ideas crystallize, collect them here before building the deck.

#	Topic	Type	Content idea	Source/visual	Status
1	Regulatory Landscape	Table	"The Spectrum" — 7 approaches from ban to nuanced distinction	See Section I notes	Draft ready
2	Regulatory Landscape	Bullets	"Trends" — 5 bullet points on direction of travel	See Section I notes	Draft ready

#	Topic	Type	Content idea	Source/visual	Status
3	Regulatory Landscape	Bullets	"Contrasting State Approaches" — IL / NY / CA as three poles	See Section I notes	Draft ready
4	Hallucination Crisis	Bullets	"The Scale" — 6 escalating stats from Charlotin database	See Section II-A notes	Draft ready
5	Hallucination Crisis	Table	"Escalating Consequences" — 7 cases from $5K to default judgment	See Section II-B notes	Draft ready
6	Hallucination Crisis	Bullets	"The Cover-Up Is Worse Than the Crime" — Cassata/Kjoller/Hayes + emerging rule	See Section II-E notes	Draft ready
7	Hallucination Crisis	Bullets	"The Detection Problem" — why detection can't solve this + arms race + pivot to verification	See Section II-B2 notes	Draft ready
8	Hallucination Crisis	Bullets	"Beyond Fake Citations" — blurring, disregarding sources, why these are worse	See Section II-C notes	Draft ready
9					
10					
11					
12					

Session Log

Brief notes after each working session so we pick up where we left off.

Session 1 — Feb 14, 2026

- Created working notes file
- Developed 8 slide drafts:
  - Slides 1-3: Regulatory Landscape (Spectrum table, Trends, Contrasting State Approaches)
  - Slides 4-8: Hallucination Crisis (The Scale, Escalating Consequences, Cover-Up Pattern + audience talking point, Detection Problem, Beyond Fake Citations)
- Corrected Charlotin database count: 851 cases, not ~100
- Discussed Stanford hallucination rate study — 2024 data, decided to date-stamp it fairly
- Refined Kjoller bullet for clarity
- Developed audience-facing talking point: "not always who you think" — good lawyers getting tripped up, not just bad actors
- **Next: Professional Responsibility (Block II) — Rules, supervision, legislative turn, disclosure**